

Async Rust: compartiendo memoria

Metodos para manejar estado mutable compartido en Async Rust


27/11/2025

Gabriel A. Steinberg




1 Un server TCP

Protocolo




- El servidor le asigna un ID a cada cliente
- Los clientes se informan su ID por un canal fuera de banda
- Cada cliente le puede pedir al servidor que envíe un mensaje

Secuencia




Async

- Esperar eventos concurrentemente
- Ejecutar tareas paralelamente



El server necesita concurrencia

Un modelo mental útil



2 Implementacion

Un Server

```
1 struct Server {  
2     listener: tokio::net::TcpListener,  
3     connections: HashMap<u32, tokio::net::TcpStream>,  
4 }
```

Los tipos pueden llegar a cambiar pero se mantiene esta idea.

Importante: Estamos usando la version async (de tokio) de `TcpListener` y `TcpStream`

Implementacion Naive (no compila)

```
1 impl Server {
2     pub async fn handle_connections(&mut self) {
3         loop {
4             let (socket, _) = self.listener.accept().await.unwrap();
5
6             tokio::spawn(async move {
7                 self.handle_connection(socket).await;
8             });
9         }
10    }
11
12    async fn handle_connection(&mut self, mut socket: TcpStream) {...}
13 }
```

El error

```
1 impl Server {  
2     pub async fn handle_connections(&mut self) {  
3         loop {  
4             let (socket, _) = self.listener.accept().await.unwrap();  
5  
6             tokio::spawn(async move {  
7                 self.handle_connection(socket).await;  
8             });  
9         }  
10    }  
11  
12    async fn handle_connection(&mut self, mut socket: TcpStream) {...}  
13 }
```

borrow escapes function

Spawn

```
1 pub fn spawn<F>(future: F) -> JoinHandle<F::Output> where
2     F: Future + Send + 'static,
3     F::Output: Send + 'static,
```

- El `Future` a ejecutar tiene que ser `'static` porque va a vivir mas que la funcion donde se llama.
- El borrow mutable a `self` continua viviendo despues de que termina la iteracion del `loop` por lo que rompe las reglas del borrow checker, si queremos mantener concurrencia.

Conclusion: Necesitamos algun mecanismo de sincronizacion.

Implementacion con Mutex

Vamos a usar un `tokio::sync::Mutex` porque vamos a necesitar tener el lock mientras llamamos `await`.

```
1 async fn handle_connection(&self, mut socket: TcpStream) {  
2     ...  
3     let (dest, msg) = parse_message(&mut buffer);  
4     self.connections.lock().get(dest).send(msg).await;  
5     ...  
6 }
```

Nuevo server

```
1 struct Server {  
2     listener: TcpListener,  
3     connections: tokio::sync::Mutex<HashMap<u32, TcpStream>>,  
4 }
```

```
1 pub async fn handle_connections(self: Arc<Self>) {  
2     loop {  
3         let (socket, _) = self.listener.accept().await.unwrap();  
4         let server = Arc::clone(&self);  
5  
6         tokio::spawn(async move { server.handle_connection(socket).await; });  
7     }  
8 }
```

Handle Connections (Mutex)

```
1  async fn handle_connection(&self, mut socket: TcpStream) {  
2      ...  
3      self.connections.insert(id, socket);  
4      loop {  
5          let connection = self.connections.lock().await.get_mut(&id).unwrap();  
6          connection.read_buf(&mut buffer).await  
7  
8          let Ok((dest, m)) = parse_message(&mut buffer) else { continue; };  
9  
10         let mut connections = self.connections.lock().await;  
11  
12         connections.get_mut(&dest).write_all(&m).await;  
13     }  
14 }
```

Deadlock

```
1  async fn handle_connection(&self, mut socket: TcpStream) {  
2      ...  
3      self.connections.insert(id, socket);  
4      loop {  
5          let connection = self.connections.lock().await.get_mut(&id).unwrap();  
6          connection.read_buf(&mut buffer).await  
7  
8          let Ok((dest, m)) = parse_message(&mut buffer) else { continue; };  
9  
10         let mut connections = self.connections.lock().await;  
11  
12         connections.get_mut(&dest).write_all(&m).await;  
13     }  
14 }
```

La solucion

```
1 struct Server {  
2     listener: TcpListener,  
3     connections: tokio::sync::Mutex<HashMap<u32, OwnedWriteHalf>>,  
4 }
```

Handle Connection

```
1  async fn handle_connection(&self, mut socket: TcpStream) {  
2      ...  
3      let (mut read, mut write) = socket.into_split();  
4      self.connections.insert(id, write);  
5      loop {  
6          let connection = read.await.get_mut(&id).unwrap();  
7          connection.read_buf(&mut buffer).await  
8  
9          let Ok((dest, m)) = parse_message(&mut buffer) else { continue; };  
10  
11         let mut connections = self.connections.lock().await;  
12         connections.get_mut(&dest).write_all(&m).await;  
13     }  
14 }
```

Aun asi...

```
1  async fn handle_connection(&self, mut socket: TcpStream) {  
2      ...  
3      let (mut read, mut write) = socket.into_split();  
4      self.connections.insert(id, write);  
5      loop {  
6          let connection = read.await.get_mut(&id).unwrap();  
7          connection.read_buf(&mut buffer).await  
8  
9          let Ok((dest, m)) = parse_message(&mut buffer) else { continue; };  
10  
11         let mut connections = self.connections.lock().await;  
12         connections.get_mut(&dest).write_all(&m).await;  
13     }  
14 }
```


Los mutex son difíciles :(

- Un buffer lleno previene a cualquier otro cliente avanzar
- Podemos usar `Arc<Mutex<OwnedWriteHalf>>` pero es mas complejo aun!

Veamos otra manera

3 Canales

Idea general



Implementacion

```
1  async fn handle_connection(&self, mut socket: TcpStream) {  
2      ...  
3      let (mut read, mut write) = socket.into_split();  
4      self.tx.send(Message::new(id, write)).await;  
5  
6      loop {  
7          let connection = read.await.get_mut(&id).unwrap();  
8          connection.read_buf(&mut buffer).await  
9  
10         let Ok((dest, m)) = parse_message(&mut buffer) else { continue; };  
11         self.tx.send(Message::Send(dest, m)).await;  
12     }  
13 }
```

Message dispatcher

```
1  async fn message_dispatcher(mut rx: mpsc::Receiver<Message>) {  
2      let mut connections = HashMap::new();  
3      while let Some(msg) = rx.recv().await {  
4          match msg {  
5              Message::New(id, mut connection) => {  
6                  connection.write_u32(id).await;  
7                  connections.insert(id, connection);  
8              }  
9              Message::Send(id, items) => {  
10                  connections.get_mut(&id).write_all(&items).await;  
11              }  
12          }  
13      }  
14  }
```

Aun asi puede bloquear...

```
1  async fn message_dispatcher(mut rx: Receiver<Message>) {  
2      let mut connections = HashMap::new();  
3      while let Some(msg) = rx.recv().await {  
4          match msg {  
5              Message::New(id, mut connection) => {  
6                  connection.write_u32(id).await;  
7                  connections.insert(id, connection);  
8              }  
9              Message::Send(id, items) => {  
10                  connections.get_mut(&id).write_all(&items).await;  
11              }  
12          }  
13      }  
14  }
```

Una mejora

```
1  async fn message_dispatcher(mut rx: Receiver<Message>) {
2      let mut connections = HashMap::new();
3      while let Some(msg) = rx.recv().await {
4          match msg {
5              Message::New(id, mut connection) => {
6                  let (tx, rx) = mpsc::channel(100);
7                  tokio::spawn(client_dispatcher(connection, rx))
8                  ...
9              }
10             Message::Send(id, items) => {
11                 connections.get_mut(&id).send(items).await;
12             }
13         }
14     }}
```

Una mejora (ii)

```
1 async fn client_dispatcher(mut connection: OwnedWriteHalf, mut rx: Receiver<Bytes>) {  
2     loop {  
3         let m = rx.recv().await.unwrap();  
4         if let Err(e) = connection.write_all(&m).await {  
5             eprintln!("Failed to write to socket: {e}");  
6             break;  
7         }  
8     }  
9 }
```

Backpressure

```
1  async fn central_dispatcher(mut rx: Receiver<Message>) {  
2      ...  
3      Message::Send(id, items, permission_token) => {  
4          ...  
5          connection.send(items).await;  
6          let _ = permission_token.send(());  
7      }  
8      ...  
9 }
```

Backpressure (ii)

```
1  async fn handle_connection(&self, socket: TcpStream) {  
2      ...  
3  
4      loop {  
5          read.read_buf(&mut buffer).await  
6  
7          ...  
8  
9          let (permission_token, permission_listener) = oneshot::channel();  
10         self.tx.send(Message::Send(dest, m, permission_token)).await.unwrap();  
11         permission_token.recv().await;  
12     }  
13 }
```

Y errores?

```
1  async fn central_dispatcher(error_tx: Sender<Message>, mut rx: Receiver<Message>) {  
2      ...  
3      while let Some(msg) = rx.recv().await {  
4          match msg {  
5              Message::New(id, connection) => {  
6                  let (tx, rx) = mpsc::channel(100);  
7                  tokio::spawn(client_dispatcher(connection, rx, error_tx.clone()));  
8                  ...  
9              }  
10             Message::Error(e, src) => {  
11                 ...  
12                 connection.get(&src).send(error).await;  
13             }  
14         }  
15     }  
16 }
```

Y errores? (ii)

```
1  async fn client_dispatcher(
2      mut connection: OwnedWriteHalf,
3      mut rx: mpsc::Receiver<(Bytes, Option<u32>)>,
4      errors_tx: mpsc::Sender<Message>,
5  ) {
6      loop {
7          let (m, src) = rx.recv().await.unwrap();
8          if let Err(e) = connection.write_all(&m).await {
9              let _ = errors_tx.send(Message::Error(e, src)).await;
10             break;
11         }
12     }
13 }
```

El problema?


- La complejidad crece, con cada interaccion entre Task.
- Manejar errores no es natural necesitas agregar nuevos mecanismos de comunicacion.
- Es dificil de debuggear.

Dado que el stack trace no sigue la forma natural de function calls, rastrear de donde proviene un valor se vuelve muy complicado.

Pero podemos simplificar esto, sacrificando el paralelismo.

4 Implementacion en un solo task

Idea general



Nueva interfaz para sockets

```
1 struct ReadSocket {  
2     buffer: BytesMut,  
3     reader: OwnedReadHalf,  
4     id: u32,  
5 }
```

Nueva interfaz para sockets (ii)

```
1 impl ReadSocket {  
2     async fn read(&mut self) -> Result<Event> {  
3         loop {  
4             if let Ok(read_result) = parse_message(&mut self.buffer) {  
5                 return Ok(Event::Read(read_result));  
6             };  
7  
8             self.reader.read_buf(&mut self.buffer)  
9                 .await  
10                .map_err(|e| (e, self.id))?  
11         }  
12     }  
13 }
```

Nueva interfaz para sockets (iii)

```
1 struct WriteSocket {  
2     buffers: VecDeque<Bytes>,  
3     writer: OwnedWriteHalf,  
4     id: u32,  
5 }
```

```
1 impl WriteSocket {  
2     fn send(&mut self, buf: Bytes) {  
3         self.buffers.push_back(buf);  
4     }  
5 }
```

Nueva interfaz para sockets (iv)

```
1  impl WriteSocket {
2      async fn advance_send(&mut self) -> Result<Event> {
3          loop {
4              let Some(buffer) = self.buffers.front_mut() else {
5                  std::future::pending::<Infallible>().await;
6                  unreachable!();
7              };
8
9              self.writer.write_all_buf(buffer)
10             .await
11             .map_err(|e| (e, self.id))?;
12
13             self.buffers.pop_front();
14     }}}
```

Server

```
1 struct Server {  
2     read_connections: HashMap<u32, ReadSocket>,  
3     write_connections: HashMap<u32, WriteSocket>,  
4     listener: TcpListener,  
5 }
```

Carrera de futuros

```
1 fn next_event(&mut self) -> impl Future<Output = Result<Event>> {
2     let listen = self.listener.accept()
3         .map(|stream| Ok(Event::Connection(stream.unwrap().0)));
4
5     let reads = self.read_connections.values_mut()
6         .map(|reader| reader.read()).collect::<Vec<_>>()
7         .race();
8
9     let writes = self.write_connections.values_mut()
10        .map(|write| write.advance_send()).collect::<Vec<_>>()
11        .race();
12
13     (listen, reads, writes).race().boxed()
14 }
```

Loop principal

```
1 pub async fn handle_connections(&mut self) {  
2     loop {  
3         ...  
4             match self.next_event().await {  
5                 Event::Read((dest, items)) => {  
6                     self.write_connections.get_mut(&dest).send(items);  
7                 }  
8                 Event::Connection(tcp_stream) => {  
9                     ...  
10                    self.read_connections.insert(id, ReadSocket::new(r, id));  
11                    self.write_connections.insert(id, WriteSocket::new(w, id));  
12                }  
13            }  
14        }  
15    }  
16}
```

Backpressure

```
1  fn next_event(&mut self) -> impl Future<Output = Result<Event>> {
2      ...
3
4      if self.buffer_size() > MAX_BUFFER {
5          (listen, writes).race().boxed()
6      } else {
7          (listen, reads, writes).race().boxed()
8      }
9  }
```

Errores

```
1  pub async fn handle_connections(&mut self) {  
2      loop {  
3          let ev = match self.next_event().await {  
4              Ok(ev) => ev,  
5              Err((e, sender)) => {  
6                  self.write_connections.get(&sender).send(error);  
7                  continue;  
8              }  
9          };  
10         ...  
11     }  
12 }
```

En resumen

- Flujo natural de errores
- Acceso `&mut` al estado
- Facil de trazar el origen de los valores
- Facil de razonar

Pero podríamos tener mas?

- Necesitamos escribir los adaptadores para los combinadores
- Tenemos que multiplexar/demultiplexar los errores
- El stack trace del IO puede ser difícil de entender
- Los futuros del IO no pueden compartir acceso `&mut` a la misma memoria

5 Conclusiones

Que vimos?

- Trade-offs: Mutex vs Canales vs Race/Select
- Evaluar si necesitamos paralelismo

Consejos

- Empezar por concurrencia sin paralelismo
- En caso de necesitar paralelismo usar canales
- Mutex se necesita en casos muy específicos
- Multiples threads no siempre mejora la performance

Mas alla de esta charla

- [Mi post](#)
- [Let future be futures](#)
- [Tree structured concurrency part I](#)
- [Tree structure concurrency part II](#)
- [sans-IO: The secret to effective Rust for network services](#)

6 Dudas? Comentarios? Quejas?